

01_import_inspect

March 9, 2024

1 Flu Shot Learning: Predict H1N1 and Seasonal Flu Vaccines

The aim of this project is to predict whether people got H1N1 and seasonal flu vaccines using information they shared about their backgrounds, opinions, and health behaviours.

This is a Driven Data problem based on a telephone survey: National 2009 H1N1 Flu Survey (NHFS) in the USA.

The population was all persons 6 months or older living in the US.

To read about the background to the project, visit <https://www.drivendata.org/competitions/66/flu-shot-learning/data/>

In-depth information about the project is available here: <https://www.drivendata.org/competitions/66/flu-shot-learning/page/211/>

To get started with this project, create a Driven Data profile, enrol in this competition, and then download the data files `../data/`!

2 Import Modules

```
[23]: import sys
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
```

3 Import and Inspect Data

```
[3]: # Import Data

submission_format = pd.read_csv('../data/submission_format.csv')
training_set_features = pd.read_csv('../data/training_set_features.csv')
training_set_labels = pd.read_csv('../data/training_set_labels.csv')
test_set_features = pd.read_csv('../data/test_set_features.csv')
```

4 Submission Format

The competition requires a `.csv` file with the following:

Columns (3): `respondent_id`, `h1n1_vaccine`, `seasonal_vaccine` Rows (26709 including first row specifying column labels): 0-26707 rows of values, plus row of column labels

So, the final submission requires `float64` probabilities for each `respondent_id` for two vaccines:

- `h1n1_vaccine`
- `seasonal_vaccine`

These are the outcome variables!

In summary, we need to build a model to predict two numeric, continuous, variables which are probabilities (`float64`).

Reading the problem description, the predictions for the two target variables should be **float probabilities** ranging from 0.0 to 1.0.

The competition uses ROC AUC as its evaluation metric, which means the values we submit should be the **probabilities that a person received each vaccine**, not binary labels.

Binary labels are provided in the `training_set_labels` dataset.

Outcomes or labels are the values to predict, in this case, float probabilities.

Predictors or features are the independent variables which explain the outcomes.

In summary, this is a supervised learning problem, because the outcome variables are labelled. We need to build a model which predicts these labels, so that probability predictions can be made for new, unseen data. In this case, the test set - which we have available (the test set it seen), and the unseen test set used by DrivenData to evaluate our project.

```
[4]: !tail -10 ../data/submission_format.csv
```

```
53405,0.5,0.7
53406,0.5,0.7
53407,0.5,0.7
53408,0.5,0.7
53409,0.5,0.7
53410,0.5,0.7
53411,0.5,0.7
53412,0.5,0.7
53413,0.5,0.7
53414,0.5,0.7
```

```
[5]: !wc ../data/submission_format.csv
```

```
26709 26709 373956 ../data/submission_format.csv
```

```
[6]: submission_format
```

```
[6]:
```

	respondent_id	h1n1_vaccine	seasonal_vaccine
0	26707	0.5	0.7
1	26708	0.5	0.7
2	26709	0.5	0.7
3	26710	0.5	0.7
4	26711	0.5	0.7
...
26703	53410	0.5	0.7
26704	53411	0.5	0.7
26705	53412	0.5	0.7
26706	53413	0.5	0.7
26707	53414	0.5	0.7

[26708 rows x 3 columns]

```
[7]: # Respondent ID: int64 dtype
submission_format['respondent_id']
```

```
[7]:
```

0	26707
1	26708
2	26709
3	26710
4	26711
...	...
26703	53410
26704	53411
26705	53412
26706	53413
26707	53414

Name: respondent_id, Length: 26708, dtype: int64

```
[8]: # h1n1_vaccine: float64 dtype
submission_format['h1n1_vaccine']
```

```
[8]:
```

0	0.5
1	0.5
2	0.5
3	0.5
4	0.5
...	...
26703	0.5
26704	0.5
26705	0.5
26706	0.5
26707	0.5

Name: h1n1_vaccine, Length: 26708, dtype: float64

```
[9]: # seasonal_vaccine: float64 dtype
      submission_format['seasonal_vaccine']
```

```
[9]: 0      0.7
      1      0.7
      2      0.7
      3      0.7
      4      0.7
      ...
      26703  0.7
      26704  0.7
      26705  0.7
      26706  0.7
      26707  0.7
      Name: seasonal_vaccine, Length: 26708, dtype: float64
```

```
[10]: training_set_features
```

```
[10]:      respondent_id  h1n1_concern  h1n1_knowledge  behavioral_antiviral_meds  \
0                0              1.0              0.0                      0.0
1                1              3.0              2.0                      0.0
2                2              1.0              1.0                      0.0
3                3              1.0              1.0                      0.0
4                4              2.0              1.0                      0.0
...            ...              ...              ...                      ...
26702           26702              2.0              0.0                      0.0
26703           26703              1.0              2.0                      0.0
26704           26704              2.0              2.0                      0.0
26705           26705              1.0              1.0                      0.0
26706           26706              0.0              0.0                      0.0

      behavioral_avoidance  behavioral_face_mask  behavioral_wash_hands  \
0                0.0              0.0              0.0
1                1.0              0.0              1.0
2                1.0              0.0              0.0
3                1.0              0.0              1.0
4                1.0              0.0              1.0
...            ...              ...              ...
26702           1.0              0.0              0.0
26703           1.0              0.0              1.0
26704           1.0              1.0              1.0
26705           0.0              0.0              0.0
26706           1.0              0.0              0.0

      behavioral_large_gatherings  behavioral_outside_home  \
0                0.0              1.0
1                0.0              1.0
```

2	0.0	0.0
3	1.0	0.0
4	1.0	0.0
...
26702	0.0	1.0
26703	0.0	0.0
26704	1.0	0.0
26705	0.0	0.0
26706	0.0	0.0

	behavioral_touch_face	...	income_poverty	marital_status	\
0	1.0	...	Below Poverty	Not Married	
1	1.0	...	Below Poverty	Not Married	
2	0.0	...	<= \$75,000, Above Poverty	Not Married	
3	0.0	...	Below Poverty	Not Married	
4	1.0	...	<= \$75,000, Above Poverty	Married	
...	
26702	0.0	...	<= \$75,000, Above Poverty	Not Married	
26703	0.0	...	<= \$75,000, Above Poverty	Not Married	
26704	1.0	...	NaN	Not Married	
26705	NaN	...	<= \$75,000, Above Poverty	Married	
26706	0.0	...	<= \$75,000, Above Poverty	Married	

	rent_or_own	employment_status	hhs_geo_region	\
0	Own	Not in Labor Force	oxchjgsf	
1	Rent	Employed	bhuqouqj	
2	Own	Employed	qufhixun	
3	Rent	Not in Labor Force	lrircsnp	
4	Own	Employed	qufhixun	
...	
26702	Own	Not in Labor Force	qufhixun	
26703	Rent	Employed	lzgpxyit	
26704	Own	NaN	lzgpxyit	
26705	Rent	Employed	lrircsnp	
26706	Own	Not in Labor Force	mlyzmhmf	

	census_msa	household_adults	household_children	\
0	Non-MSA	0.0	0.0	
1	MSA, Not Principle City	0.0	0.0	
2	MSA, Not Principle City	2.0	0.0	
3	MSA, Principle City	0.0	0.0	
4	MSA, Not Principle City	1.0	0.0	
...	
26702	Non-MSA	0.0	0.0	
26703	MSA, Principle City	1.0	0.0	
26704	MSA, Not Principle City	0.0	0.0	
26705	Non-MSA	1.0	0.0	

26706	MSA, Principle City	1.0	0.0
-------	---------------------	-----	-----

	employment_industry	employment_occupation
0	NaN	NaN
1	pxcmvdjn	xgwztkwe
2	rucpziij	xtkaffoo
3	NaN	NaN
4	wxleyezf	emcorrxb
...
26702	NaN	NaN
26703	fcxhlnwr	cmhcxjea
26704	NaN	NaN
26705	fcxhlnwr	haliazsg
26706	NaN	NaN

[26707 rows x 36 columns]

```
[11]: training_set_features.columns
```

```
[11]: Index(['respondent_id', 'h1n1_concern', 'h1n1_knowledge',
          'behavioral_antiviral_meds', 'behavioral_avoidance',
          'behavioral_face_mask', 'behavioral_wash_hands',
          'behavioral_large_gatherings', 'behavioral_outside_home',
          'behavioral_touch_face', 'doctor_recc_h1n1', 'doctor_recc_seasonal',
          'chronic_med_condition', 'child_under_6_months', 'health_worker',
          'health_insurance', 'opinion_h1n1_vacc_effective', 'opinion_h1n1_risk',
          'opinion_h1n1_sick_from_vacc', 'opinion_seas_vacc_effective',
          'opinion_seas_risk', 'opinion_seas_sick_from_vacc', 'age_group',
          'education', 'race', 'sex', 'income_poverty', 'marital_status',
          'rent_or_own', 'employment_status', 'hhs_geo_region', 'census_msa',
          'household_adults', 'household_children', 'employment_industry',
          'employment_occupation'],
          dtype='object')
```

5 Training Set Features

- All the numbers, except `respondent_id` (int) are floats, but could be converted to integers (ordinal or nominal)
- There are quite a few features; these are the predictor variables
- Contains 26_707 rows, but the `submission_format` contains 26_708 rows. This is due to the 1st row being the column labels.

Reading the problem description, we discover:

- 35 potential predictors (`respondent_id` is a unique and random identifier)

5.1 Binary Variables

behavioral_antiviral_meds - Has taken antiviral medications. (binary) **behavioral_avoidance** - Has avoided close contact with others with flu-like symptoms. (binary) **behavioral_face_mask** - Has bought a face mask. (binary) **behavioral_wash_hands** - Has frequently washed hands or used hand sanitizer. (binary) **behavioral_large_gatherings** - Has reduced time at large gatherings. (binary) **behavioral_outside_home** - Has reduced contact with people outside of own household. (binary) **behavioral_touch_face** - Has avoided touching eyes, nose, or mouth. (binary) **doctor_recc_h1n1** - H1N1 flu vaccine was recommended by doctor. (binary) **doctor_recc_seasonal** - Seasonal flu vaccine was recommended by doctor. (binary) **chronic_med_condition** - Has any of the following chronic medical conditions: asthma or an other lung condition, diabetes, a heart condition, a kidney condition, sickle cell anemia or other anemia, a neurological or neuromuscular condition, a liver condition, or a weakened immune system caused by a chronic illness or by medicines taken for a chronic illness. (binary) **child_under_6_months** - Has regular close contact with a child under the age of six months. (binary) **health_worker** - Is a healthcare worker. (binary) **health_insurance** - Has health insurance. (binary)

6 Ordinal Variables (Scales)

h1n1_concern - Level of concern about the H1N1 flu.

0 = Not at all concerned; 1 = Not very concerned; 2 = Somewhat concerned; 3 = Very concerned.

h1n1_knowledge - Level of knowledge about H1N1 flu.

0 = No knowledge; 1 = A little knowledge; 2 = A lot of knowledge.

opinion_h1n1_vacc_effective - Respondent's opinion about H1N1 vaccine effectiveness.

1 = Not at all effective; 2 = Not very effective; 3 = Don't know; 4 = Somewhat effective; 5 = Very effective.

opinion_h1n1_risk - Respondent's opinion about risk of getting sick with H1N1 flu without vaccine.

1 = Very Low; 2 = Somewhat low; 3 = Don't know; 4 = Somewhat high; 5 = Very high.

opinion_h1n1_sick_from_vacc - Respondent's worry of getting sick from taking H1N1 vaccine.

1 = Not at all worried; 2 = Not very worried; 3 = Don't know; 4 = Somewhat worried; 5 = Very worried.

opinion_seas_vacc_effective - Respondent's opinion about seasonal flu vaccine effectiveness.

1 = Not at all effective; 2 = Not very effective; 3 = Don't know; 4 = Somewhat effective; 5 = Very effective.

opinion_seas_risk - Respondent's opinion about risk of getting sick with seasonal flu without vaccine.

1 = Very Low; 2 = Somewhat low; 3 = Don't know; 4 = Somewhat high; 5 = Very high.

opinion_seas_sick_from_vacc - Respondent's worry of getting sick from taking seasonal flu vaccine.

1 = Not at all worried; 2 = Not very worried; 3 = Don't know; 4 = Somewhat worried; 5 = Very worried.

7 Demographics

`age_group` - Age group of respondent. `education` - Self-reported education level. `race` - Race of respondent. `sex` - Sex of respondent.

8 Character Variables

`hhs_geo_region` - Respondent's residence using a 10-region geographic classification defined by the U.S. Dept. of Health and Human Services. Values are represented as short random character strings. `employment_industry` - Type of industry respondent is employed in. Values are represented as short random character strings. `employment_occupation` - Type of occupation of respondent. Values are represented as short random character strings.

9 Discrete Variables (Counts)

`household_adults` - Number of other adults in household, top-coded to 3. `household_children` - Number of children in household, top-coded to 3.

10 Residence

`hhs_geo_region` - Respondent's residence using a 10-region geographic classification defined by the U.S. Dept. of Health and Human Services. Values are represented as short random character strings. `census_msa` - Respondent's residence within metropolitan statistical areas (MSA) as defined by the U.S. Census.

11 Misc

`income_poverty` - Household annual income of respondent with respect to 2008 Census poverty thresholds. `marital_status` - Marital status of respondent. `rent_or_own` - Housing situation of respondent. `employment_status` - Employment status of respondent.

12 Training Set Labels

DataFrame with 26707 rows and 3 columns.

The columns are the same as the columns in the `submission_format.csv` file:

- `respondent_id`: integers (0-26_706)
- `h1n1_vaccine`: 0 (No) or 1 (Yes) (Boolean): whether respondent received H1N1 flu vaccine.
- `seasonal_vaccine`: 0 (No) or 1 (Yes) (Boolean): whether respondent received seasonal flue vaccine.

These are the outcome variables which are `True` or `False` for whether people did or did not get vaccines for H1N1 or seasonal flu.

The outcome variables, or dependent variables, are binary. This influences the kinds of exploratory data analysis and statistical modeling we can do.

DrivenData describes this as a **multilabel** - not multiclass - problem.

As this is a **multilabel** problem, the **probabilities for each row do not need to sum to one**.

```
[12]: training_set_labels
```

```
[12]:      respondent_id  h1n1_vaccine  seasonal_vaccine
0                0              0              0
1                1              0              1
2                2              0              0
3                3              0              1
4                4              0              0
...
26702           26702              0              0
26703           26703              0              0
26704           26704              0              1
26705           26705              0              0
26706           26706              0              0
```

```
[26707 rows x 3 columns]
```

13 Test Set Features

DataFrame with 26708 rows and 36 columns

- 26707 rows in the `training_set_features` DataFrame
- Eyeballing the DataFrame, the columns look the same as the `training_set_features`

This is the dataset we use to test our model predictions.

It would be good to have a holdout or validation set from the training set.

Leave this test set alone, until we have built and validated our model!

```
[13]: test_set_features
```

```
[13]:      respondent_id  h1n1_concern  h1n1_knowledge  behavioral_antiviral_meds  \
0                26707            2.0            2.0                    0.0
1                26708            1.0            1.0                    0.0
2                26709            2.0            2.0                    0.0
3                26710            1.0            1.0                    0.0
4                26711            3.0            1.0                    1.0
...
26703           53410            1.0            1.0                    0.0
26704           53411            3.0            1.0                    0.0
26705           53412            0.0            1.0                    0.0
26706           53413            3.0            1.0                    0.0
26707           53414            2.0            1.0                    0.0

      behavioral_avoidance  behavioral_face_mask  behavioral_wash_hands  \
0                1.0                0.0                1.0
```

1	0.0	0.0	0.0
2	0.0	1.0	1.0
3	0.0	0.0	0.0
4	1.0	0.0	1.0
...
26703	1.0	0.0	1.0
26704	1.0	0.0	1.0
26705	0.0	0.0	0.0
26706	1.0	0.0	1.0
26707	0.0	0.0	1.0

	behavioral_large_gatherings	behavioral_outside_home \
0	1.0	0.0
1	0.0	0.0
2	1.0	1.0
3	0.0	0.0
4	1.0	1.0
...
26703	0.0	0.0
26704	1.0	1.0
26705	0.0	0.0
26706	0.0	1.0
26707	0.0	0.0

	behavioral_touch_face	...	income_poverty	marital_status \
0	1.0	...	> \$75,000	Not Married
1	0.0	...	Below Poverty	Not Married
2	1.0	...	> \$75,000	Married
3	0.0	...	<= \$75,000, Above Poverty	Married
4	1.0	...	<= \$75,000, Above Poverty	Not Married
...
26703	1.0	...	NaN	NaN
26704	1.0	...	Below Poverty	Married
26705	0.0	...	Below Poverty	Not Married
26706	0.0	...	<= \$75,000, Above Poverty	Married
26707	1.0	...	NaN	Not Married

	rent_or_own	employment_status	hhs_geo_region \
0	Rent	Employed	mlyzmhmf
1	Rent	Employed	bhuqouqj
2	Own	Employed	lrircsnp
3	Own	Not in Labor Force	lrircsnp
4	Own	Employed	lzgpxyit
...
26703	NaN	NaN	dqpwygqj
26704	Rent	Employed	qufhixun
26705	Rent	Not in Labor Force	qufhixun

26706	Own	Not in Labor Force	bhuqouqj
26707	Rent	Employed	lrircsnp

	census_msa	household_adults	household_children	\
0	MSA, Not Principle City	1.0	0.0	
1	Non-MSA	3.0	0.0	
2	Non-MSA	1.0	0.0	
3	MSA, Not Principle City	1.0	0.0	
4	Non-MSA	0.0	1.0	
...	
26703	MSA, Principle City	1.0	1.0	
26704	Non-MSA	1.0	3.0	
26705	MSA, Not Principle City	1.0	0.0	
26706	MSA, Not Principle City	1.0	0.0	
26707	MSA, Principle City	0.0	0.0	

	employment_industry	employment_occupation
0	atmlpfrs	hfxkjkmi
1	atmlpfrs	xqwwgdyp
2	nduyfdeo	pvmttkik
3	NaN	NaN
4	fcxhlnwr	mxkfnird
...
26703	NaN	NaN
26704	fcxhlnwr	vlluhbov
26705	NaN	NaN
26706	NaN	NaN
26707	NaN	xtkaffoo

[26708 rows x 36 columns]

14 Summarise Datasets and Data Types

```
[14]: from ydata_profiling import ProfileReport
```

```
[15]: training_set_features.describe()
```

```
[15]:
```

	respondent_id	h1n1_concern	h1n1_knowledge	behavioral_antiviral_meds	\
count	26707.000000	26615.000000	26591.000000	26636.000000	
mean	13353.000000	1.618486	1.262532	0.048844	
std	7709.791156	0.910311	0.618149	0.215545	
min	0.000000	0.000000	0.000000	0.000000	
25%	6676.500000	1.000000	1.000000	0.000000	
50%	13353.000000	2.000000	1.000000	0.000000	
75%	20029.500000	2.000000	2.000000	0.000000	
max	26706.000000	3.000000	2.000000	1.000000	

	behavioral_avoidance	behavioral_face_mask	behavioral_wash_hands \
count	26499.000000	26688.000000	26665.000000
mean	0.725612	0.068982	0.825614
std	0.446214	0.253429	0.379448
min	0.000000	0.000000	0.000000
25%	0.000000	0.000000	1.000000
50%	1.000000	0.000000	1.000000
75%	1.000000	0.000000	1.000000
max	1.000000	1.000000	1.000000

	behavioral_large_gatherings	behavioral_outside_home \
count	26620.00000	26625.000000
mean	0.35864	0.337315
std	0.47961	0.472802
min	0.00000	0.000000
25%	0.00000	0.000000
50%	0.00000	0.000000
75%	1.00000	1.000000
max	1.00000	1.000000

	behavioral_touch_face ...	health_worker	health_insurance \
count	26579.000000 ...	25903.000000	14433.00000
mean	0.677264 ...	0.111918	0.87972
std	0.467531 ...	0.315271	0.32530
min	0.000000 ...	0.000000	0.00000
25%	0.000000 ...	0.000000	1.00000
50%	1.000000 ...	0.000000	1.00000
75%	1.000000 ...	0.000000	1.00000
max	1.000000 ...	1.000000	1.00000

	opinion_h1n1_vacc_effective	opinion_h1n1_risk \
count	26316.000000	26319.000000
mean	3.850623	2.342566
std	1.007436	1.285539
min	1.000000	1.000000
25%	3.000000	1.000000
50%	4.000000	2.000000
75%	5.000000	4.000000
max	5.000000	5.000000

	opinion_h1n1_sick_from_vacc	opinion_seas_vacc_effective \
count	26312.000000	26245.000000
mean	2.357670	4.025986
std	1.362766	1.086565
min	1.000000	1.000000
25%	1.000000	4.000000

50%	2.000000	4.000000
75%	4.000000	5.000000
max	5.000000	5.000000

	opinion_seas_risk	opinion_seas_sick_from_vacc	household_adults \
count	26193.000000	26170.000000	26458.000000
mean	2.719162	2.118112	0.886499
std	1.385055	1.332950	0.753422
min	1.000000	1.000000	0.000000
25%	2.000000	1.000000	0.000000
50%	2.000000	2.000000	1.000000
75%	4.000000	4.000000	1.000000
max	5.000000	5.000000	3.000000

	household_children
count	26458.000000
mean	0.534583
std	0.928173
min	0.000000
25%	0.000000
50%	0.000000
75%	1.000000
max	3.000000

[8 rows x 24 columns]

```
[16]: profile = ProfileReport(training_set_features, title="Profiling Report")
profile
```

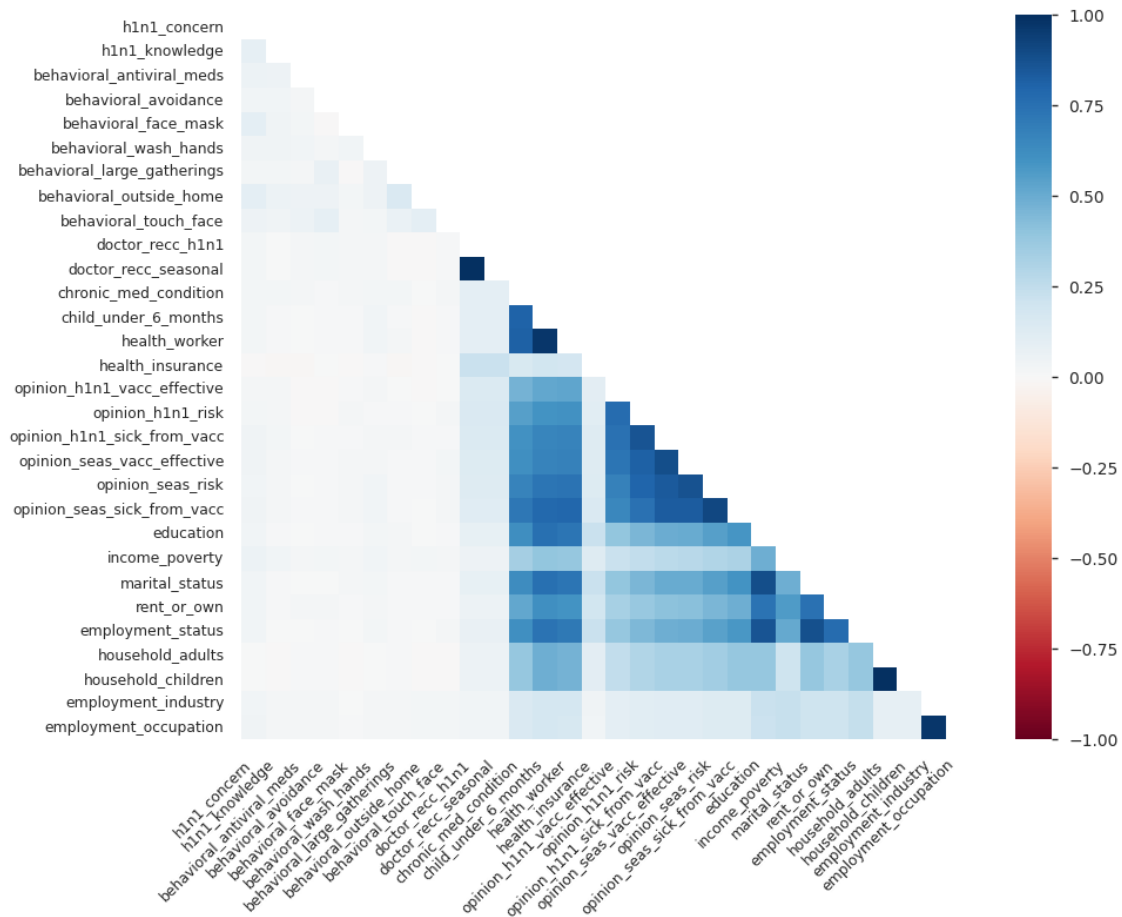
Summarize dataset: 0%| | 0/5 [00:00<?, ?it/s]

Generate report structure: 0%| | 0/1 [00:00<?, ?it/s]

Render HTML: 0%| | 0/1 [00:00<?, ?it/s]

<IPython.core.display.HTML object>

[16]:



```
[17]: training_set_labels.describe()
```

```
[17]:
```

	respondent_id	h1n1_vaccine	seasonal_vaccine
count	26707.000000	26707.000000	26707.000000
mean	13353.000000	0.212454	0.465608
std	7709.791156	0.409052	0.498825
min	0.000000	0.000000	0.000000
25%	6676.500000	0.000000	0.000000
50%	13353.000000	0.000000	0.000000
75%	20029.500000	0.000000	1.000000
max	26706.000000	1.000000	1.000000

```
[18]: profile = ProfileReport(training_set_labels, title="Profiling Report")
profile
```

```
Summarize dataset: 0%| | 0/5 [00:00<?, ?it/s]
```

```
Generate report structure: 0%| | 0/1 [00:00<?, ?it/s]
```

```
Render HTML: 0%| | 0/1 [00:00<?, ?it/s]
```

<IPython.core.display.HTML object>

[18]:

[19]: test_set_features.describe()

[19]:

	respondent_id	h1n1_concern	h1n1_knowledge	behavioral_antiviral_meds	\
count	26708.000000	26623.000000	26586.000000	26629.000000	
mean	40060.500000	1.623145	1.266042	0.049645	
std	7710.079831	0.902755	0.615617	0.217215	
min	26707.000000	0.000000	0.000000	0.000000	
25%	33383.750000	1.000000	1.000000	0.000000	
50%	40060.500000	2.000000	1.000000	0.000000	
75%	46737.250000	2.000000	2.000000	0.000000	
max	53414.000000	3.000000	2.000000	1.000000	

	behavioral_avoidance	behavioral_face_mask	behavioral_wash_hands	\
count	26495.000000	26689.000000	26668.000000	
mean	0.729798	0.069279	0.826084	
std	0.444072	0.253934	0.379045	
min	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	1.000000	
50%	1.000000	0.000000	1.000000	
75%	1.000000	0.000000	1.000000	
max	1.000000	1.000000	1.000000	

	behavioral_large_gatherings	behavioral_outside_home	\
count	26636.000000	26626.000000	
mean	0.351517	0.337227	
std	0.477453	0.472772	
min	0.000000	0.000000	
25%	0.000000	0.000000	
50%	0.000000	0.000000	
75%	1.000000	1.000000	
max	1.000000	1.000000	

	behavioral_touch_face	...	health_worker	health_insurance	\
count	26580.000000	...	25919.000000	14480.000000	
mean	0.683747	...	0.111501	0.887914	
std	0.465022	...	0.314758	0.315483	
min	0.000000	...	0.000000	0.000000	
25%	0.000000	...	0.000000	1.000000	
50%	1.000000	...	0.000000	1.000000	
75%	1.000000	...	0.000000	1.000000	
max	1.000000	...	1.000000	1.000000	

	opinion_h1n1_vacc_effective	opinion_h1n1_risk	\
count	26310.000000	26328.000000	

mean	3.844622	2.326838
std	1.007570	1.275636
min	1.000000	1.000000
25%	3.000000	1.000000
50%	4.000000	2.000000
75%	5.000000	4.000000
max	5.000000	5.000000

	opinion_h1n1_sick_from_vacc	opinion_seas_vacc_effective \
count	26333.000000	26256.000000
mean	2.360612	4.024832
std	1.359413	1.083204
min	1.000000	1.000000
25%	1.000000	4.000000
50%	2.000000	4.000000
75%	4.000000	5.000000
max	5.000000	5.000000

	opinion_seas_risk	opinion_seas_sick_from_vacc	household_adults \
count	26209.000000	26187.000000	26483.000000
mean	2.708688	2.143392	0.894310
std	1.376045	1.339102	0.754244
min	1.000000	1.000000	0.000000
25%	2.000000	1.000000	0.000000
50%	2.000000	2.000000	1.000000
75%	4.000000	4.000000	1.000000
max	5.000000	5.000000	3.000000

	household_children
count	26483.000000
mean	0.543745
std	0.935057
min	0.000000
25%	0.000000
50%	0.000000
75%	1.000000
max	3.000000

[8 rows x 24 columns]

```
[20]: profile = ProfileReport(test_set_features, title="Profiling Report")
profile
```


Summarize dataset: 0%| | 0/5 [00:00<?, ?it/s]

Generate report structure: 0%| | 0/1 [00:00<?, ?it/s]

Render HTML: 0%| | 0/1 [00:00<?, ?it/s]

<IPython.core.display.HTML object>

[20]:



15 Missing Values

I discovered the ordinal/categorical variables could not be converted to integers, due to the presence of NaNs

So, then I investigated NaNs in the data

16 Check NaN (Not A Number) Values

```
[21]: nan_count_per_column = training_set_features.isna().sum()

total_nan_count = training_set_features.isna().sum().sum()
print(total_nan_count)
```

17 Plot Missing Values

```
[24]: # Create a boolean DataFrame of missing values for all columns except
      ↪ 'respondent_id'
missing_values = training_set_features.drop(columns=['respondent_id']).isna()

# Calculate the sum of missing values for each variable and sort them in
      ↪ ascending order
sorted_columns = missing_values.sum().sort_values().index

# Reorder the missing_values DataFrame based on the sorted variables
sorted_missing_values = missing_values[sorted_columns]

# Create the figure and axis for the heatmap
plt.figure(figsize=(12, 10)) # Adjust figure size as needed
ax = sns.heatmap(sorted_missing_values.T, cbar=False, xticklabels=False,
      ↪ cmap='viridis')

# Set the background color of the figure (outside the heatmap)
plt.gcf().set_facecolor('black')

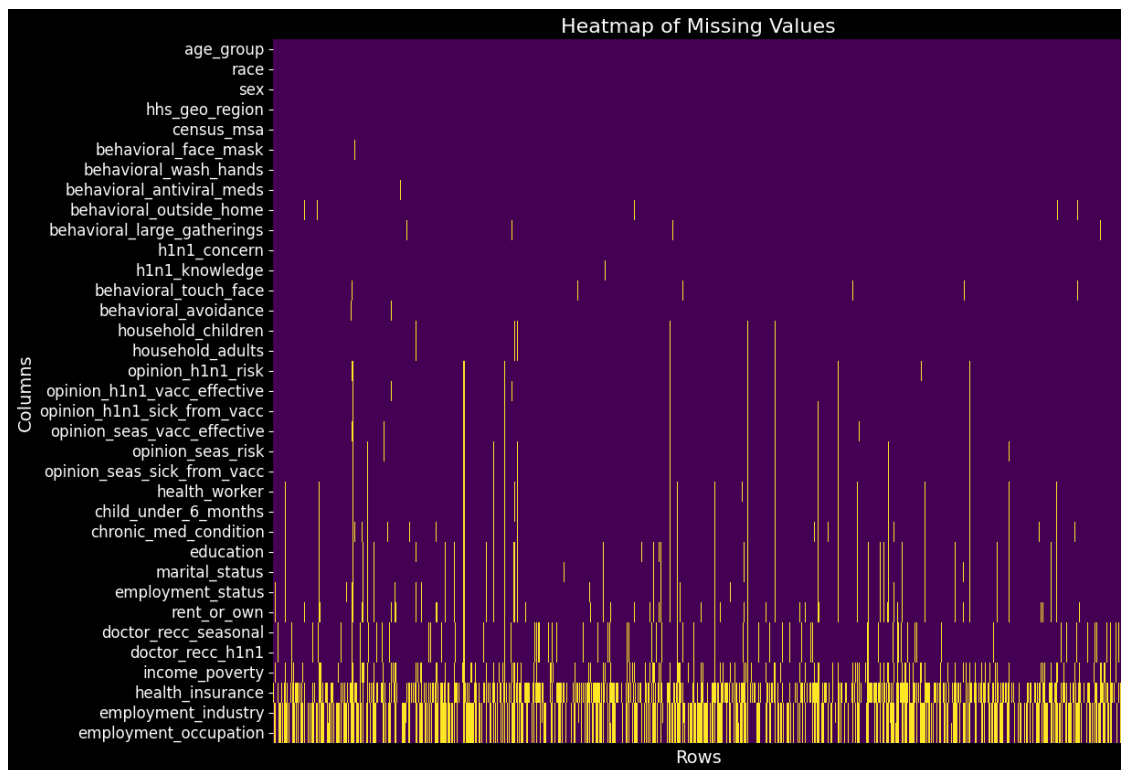
# Set the title and labels with white text
plt.title('Heatmap of Missing Values', color='white', fontsize=16)
plt.xlabel('Rows', color='white', fontsize=14)
plt.ylabel('Columns', color='white', fontsize=14)

# Set tick colors to white and increase font size for better readability
ax.tick_params(axis='y', colors='white', labelsiz=12)

# Set the color of the axis (spine) to white
for spine in ax.spines.values():
    spine.set_edgecolor('white')

# Optionally, rotate the y-axis labels for better readability if needed
plt.yticks(rotation=0)

# Save the figure with a black background
plt.savefig('../images/nan_heatmap.png', facecolor='black',
      ↪ bbox_inches='tight', pad_inches=1.0)
plt.show()
```



18 Check Missing Values

This heatmap shows the presence of NaN values in the dataset according to column (feature) variables

I was about to drop all rows with missing values (NaN)

But after plotting, I can see the missing values group around the following variables:

- **employment_occupation** has 13470 (50.4%) missing values
- **employment_industry** has 13330 (49.9%) missing values
- **health_insurance** has 12274 (46.0%) missing values
- **income_poverty** has 4423 (16.6%) missing values

This possibly suggests participants felt uncomfortable answering these questions during the telephone survey

But we don't know!

I will keep the NaNs for now.

We need to be careful with those variables where there are many NaNs

They may give indications about the nature of the data, and give us caution about how we interpret this dataset

19 Interactive Heatmap

```
[28]: import plotly.express as px

# Assuming numeric_missing_values and sorted_columns are already defined

# Increase the figure's height to give more space for y-axis labels
# The height might need further adjustment based on the actual number of labels
fig_height = max(600, 30 * len(sorted_columns))

fig = px.imshow(numeric_missing_values.T, color_continuous_scale='Viridis',
                labels=dict(x="Rows", y="Columns", color="Missing Values"),
                title="Heatmap of Missing Values",
                height=fig_height) # Set custom height

# Update layout to improve aesthetics and readability of y-axis labels
fig.update_layout(
    plot_bgcolor='black',
    paper_bgcolor='black',
    title_font=dict(size=16, color='white'),
    xaxis=dict(showticklabels=False),
    yaxis=dict(
        tickmode='array',
        tickvals=list(range(len(sorted_columns))),
        ticktext=sorted_columns,
        tickfont=dict(size=10, color='white') # Adjust font size as needed
    ),
    yaxis_title="Columns",
    xaxis_title="Rows"
)

# Show the figure
fig.show()
```

20 Evaluation Metric

This competition uses ROC AUC to evaluate performance.

ROC AUC is Receiver Operating Characteristic Curve.

The ROC AUC should be provided for each of the target variables - `h1n1_vaccine` and `seasonal_vaccine`.

The mean of these two scores will be the overall score. A higher value indicates stronger performance.

We're advised to use `sklearn.metrics.roc_auc_score` for this multi-label problem, using the default `average=macro` parameter.

The primary evaluation metric is Area Under the Receiver Operating Characteristic Curve (AUROC or AUC).

This metric is calculated for each label in the submission and then averaged across the labels.

AUCROC ranges from 0 to 1. The goal is to maximize AUROC.

$$AUROC = \int_{-\infty}^{\infty} TPR(T)FPR'(T) dT$$

For more information on AUCROC:

https://en.wikipedia.org/wiki/Receiver_operating_characteristic
[learn.org/stable/modules/generated/sklearn.metrics.roc_auc_score.html](https://scikit-learn.org/stable/modules/generated/sklearn.metrics.roc_auc_score.html)

[https://scikit-](https://scikit-learn.org/stable/modules/generated/sklearn.metrics.roc_auc_score.html)

21 TODO

21.1 Essential

- Create a validation or holdout set from the training set
- Find questions asked on the survey
- Check data values and data types
- Convert and transform data
- Build quickest model and submit result
- Revise ROC AUC evaluation metric using help.
- Clarify if the ROC AUC values are included in the submission `.csv` file, or are these values to calculate the performance of our model (and we should calculate).
- Import and inspect data
- Exploratory Data Analysis, starting with features, to notice any class imbalance, which would influence sampling from training set for validation (holdout) set.
- Build baseline model: simple as possible and good enough!
- Submit result
- Find NaN or missing or extreme/outlier values.

21.2 Nice to have

- Streamlit
- GitHub repository